

Fast garbage collection

part 1

Initial memory
layout:

1 object array
with 5 elements

Stack

array[]
Return address
MeasureA
<no params>

Heap

array[]
[0]
[1]
[2]
[3]
[4]
object 0
object 1
object 2
object 3
object 4

Stack

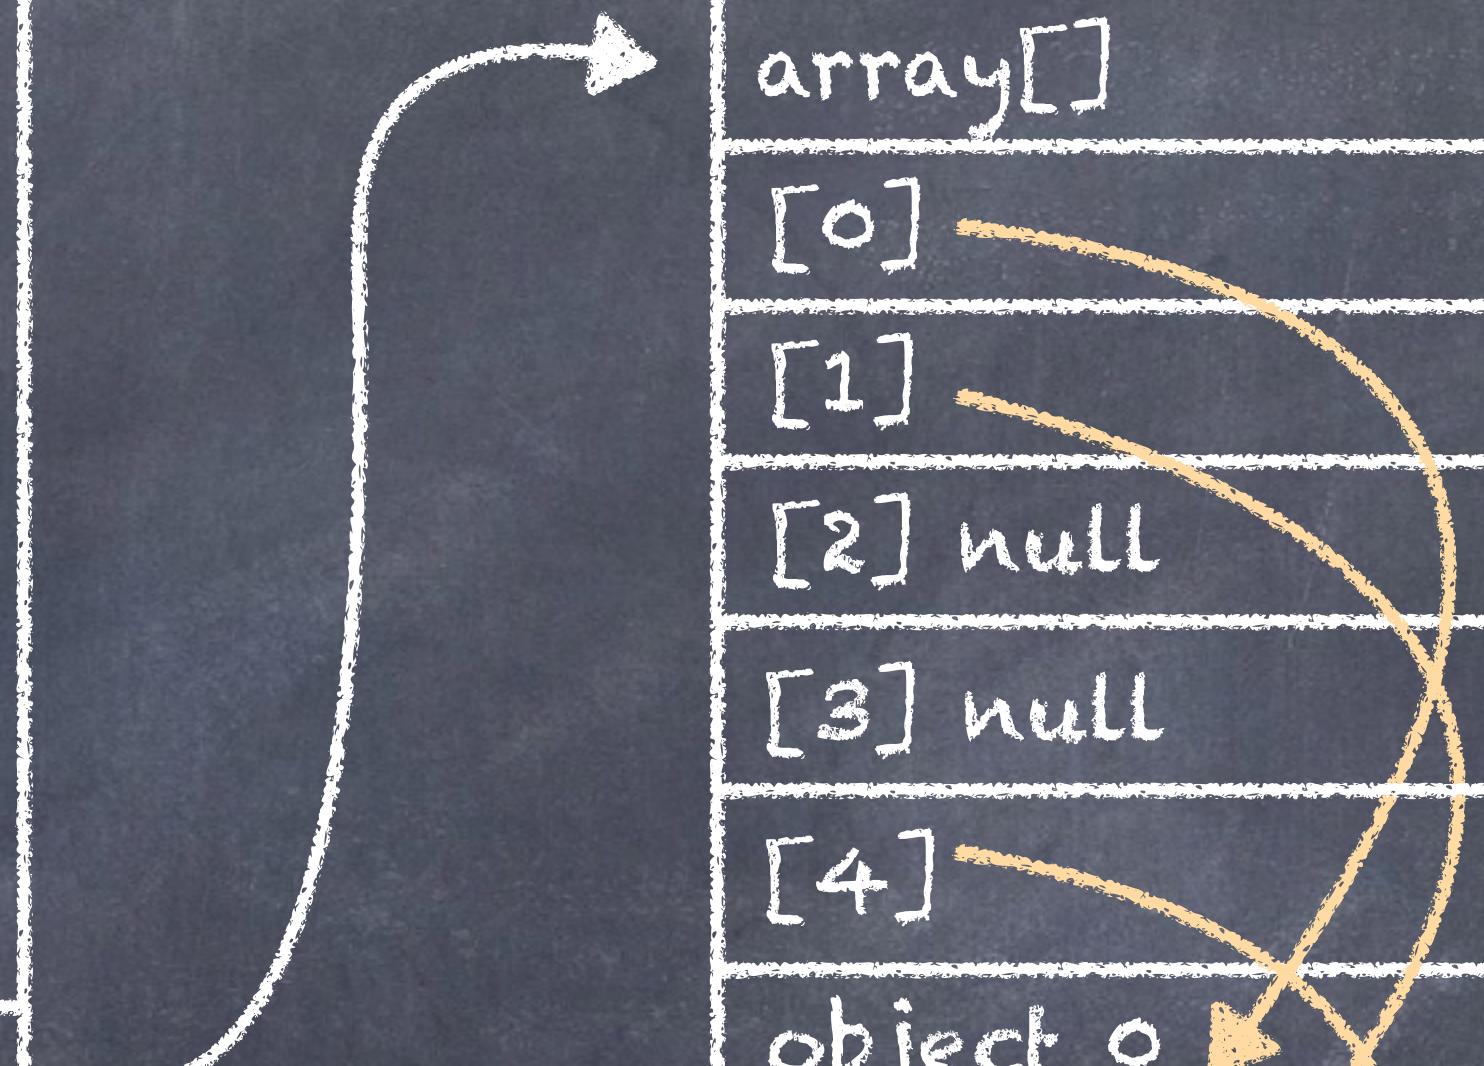
Now we set array elements 2 and 3 to null.

Objects 2 and 3 are dereferenced

array[]
Return address
MeasureA
<no params>

Heap

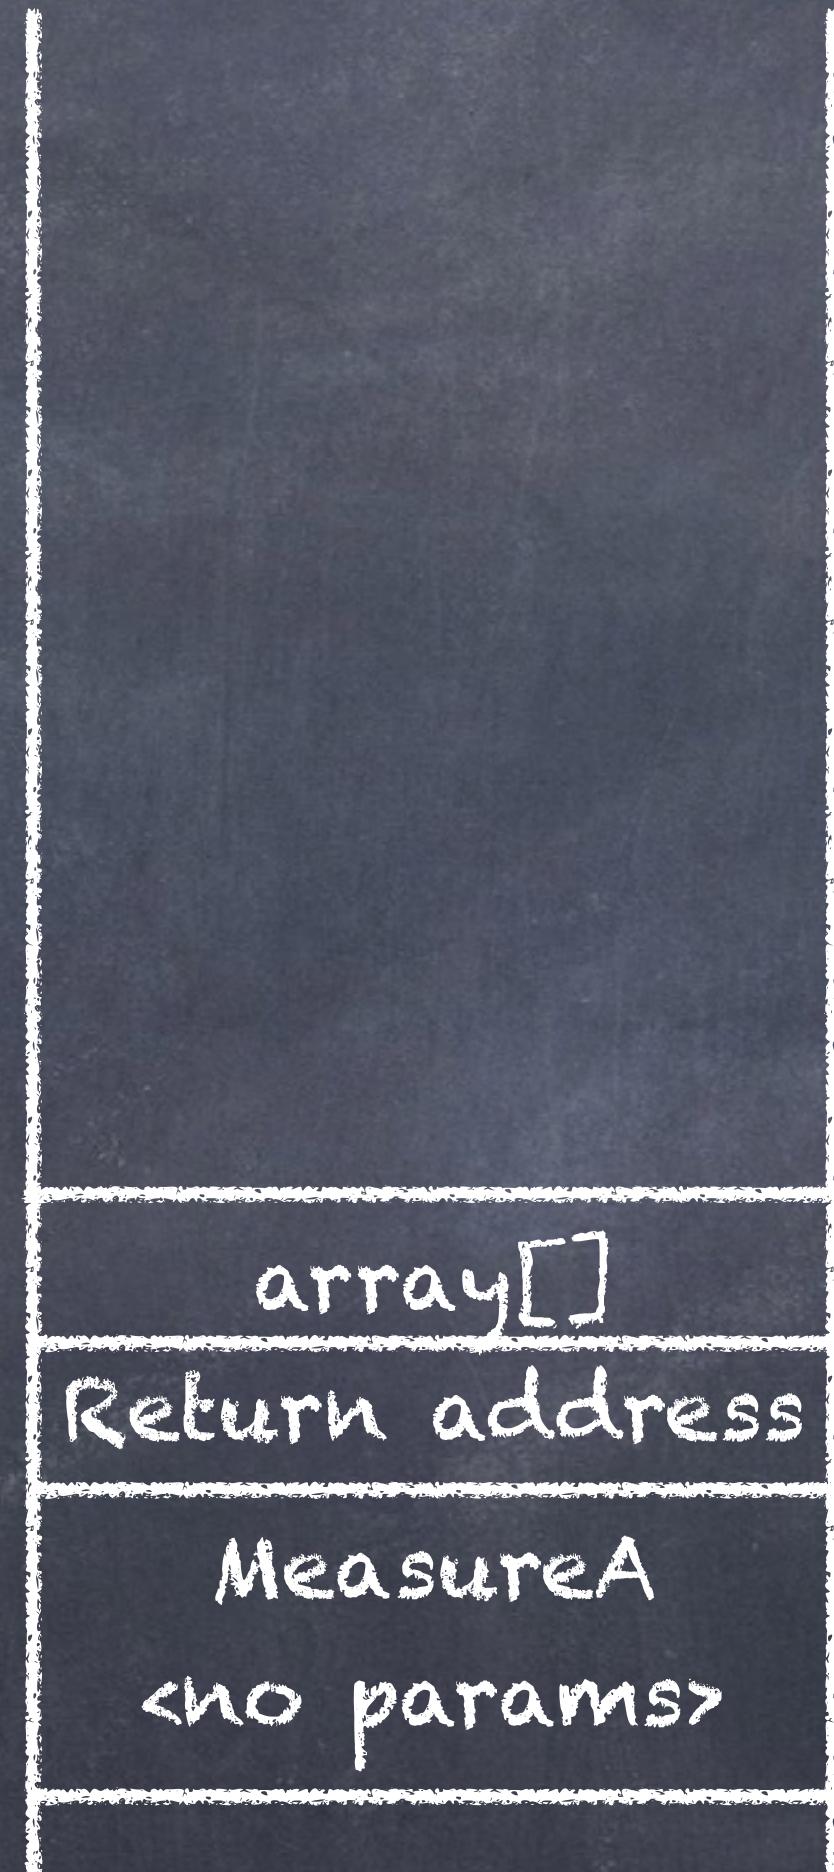
array[]
[0]
[1]
[2] null
[3] null
[4]
object 0
object 1
object 2
object 3
object 4



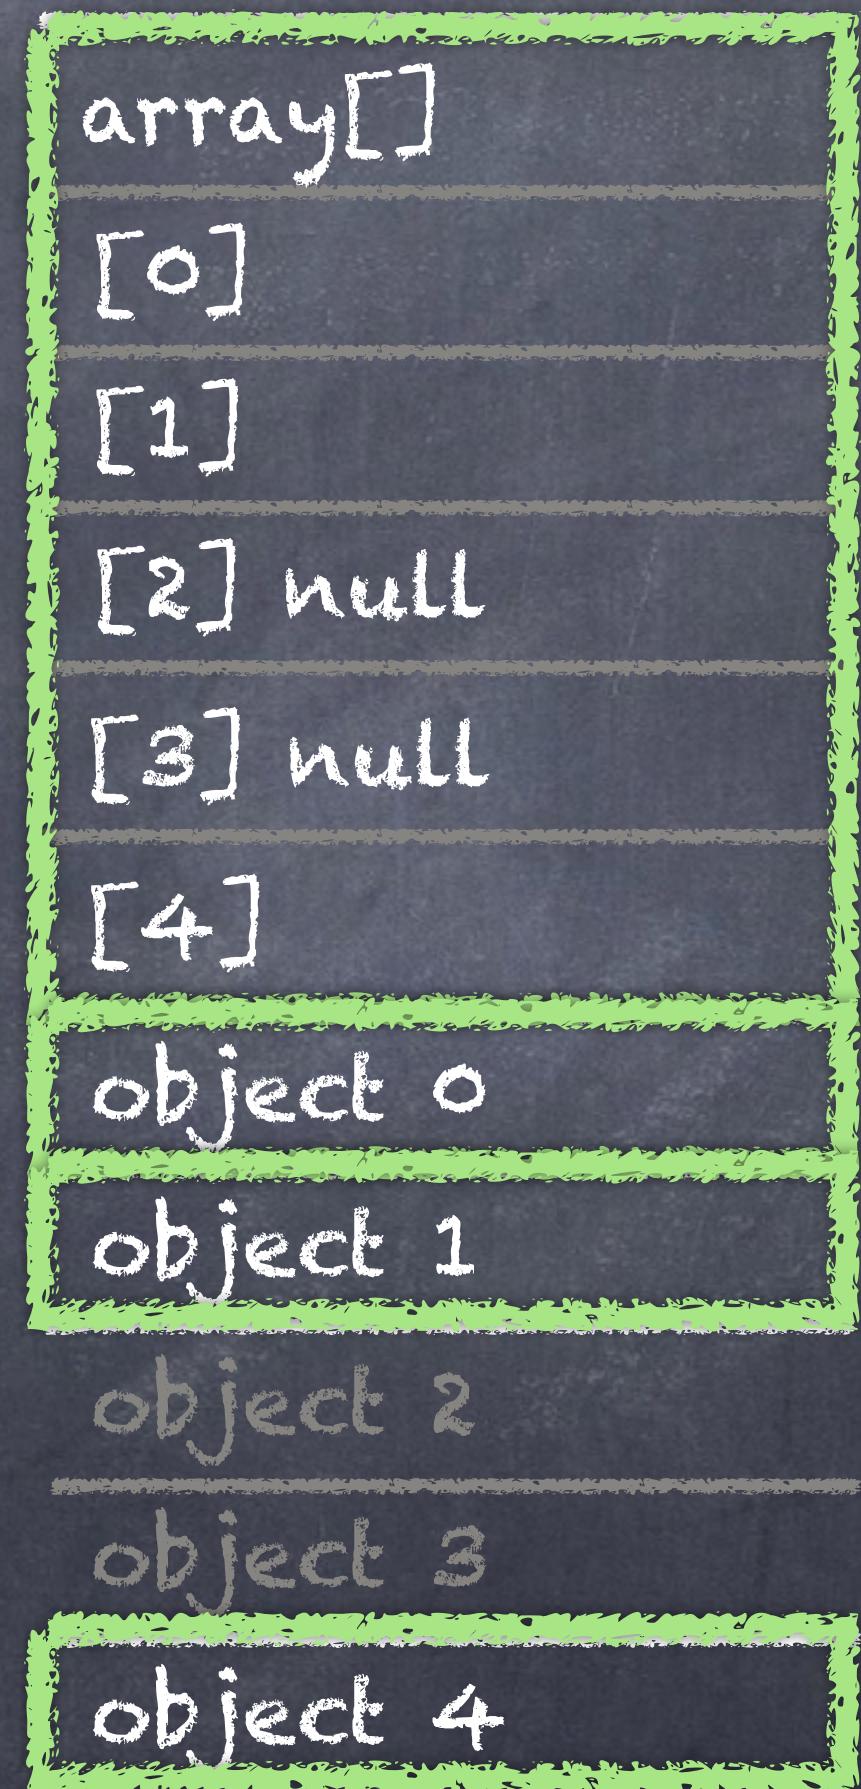
Garbage
collection runs:

All live objects
are marked

Stack



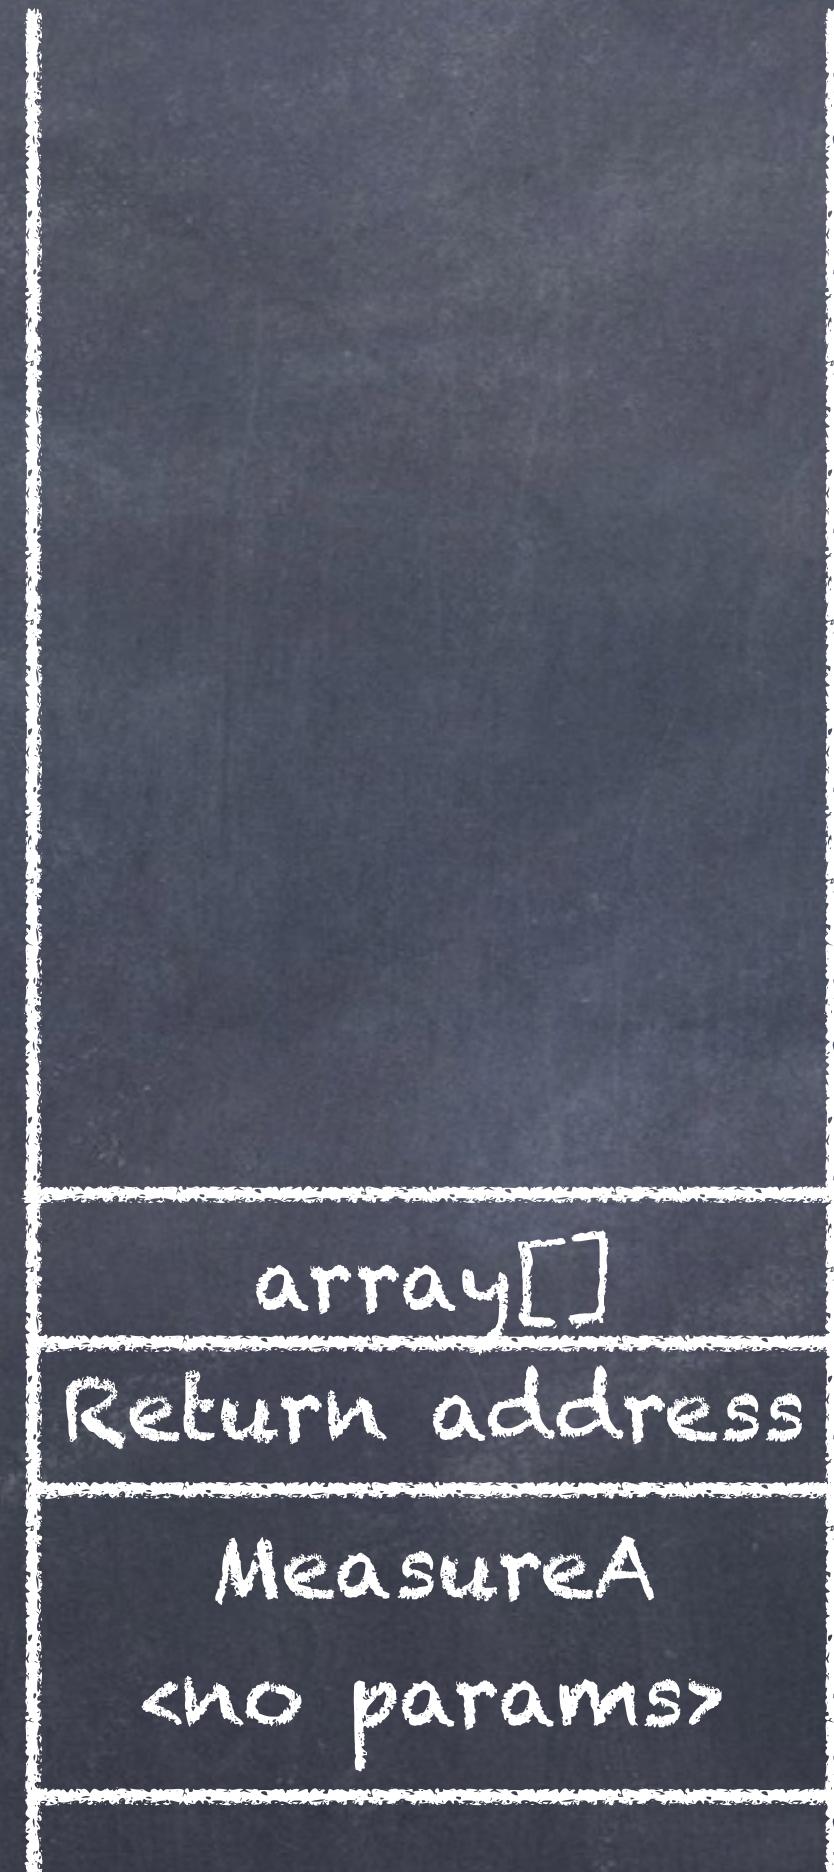
Heap



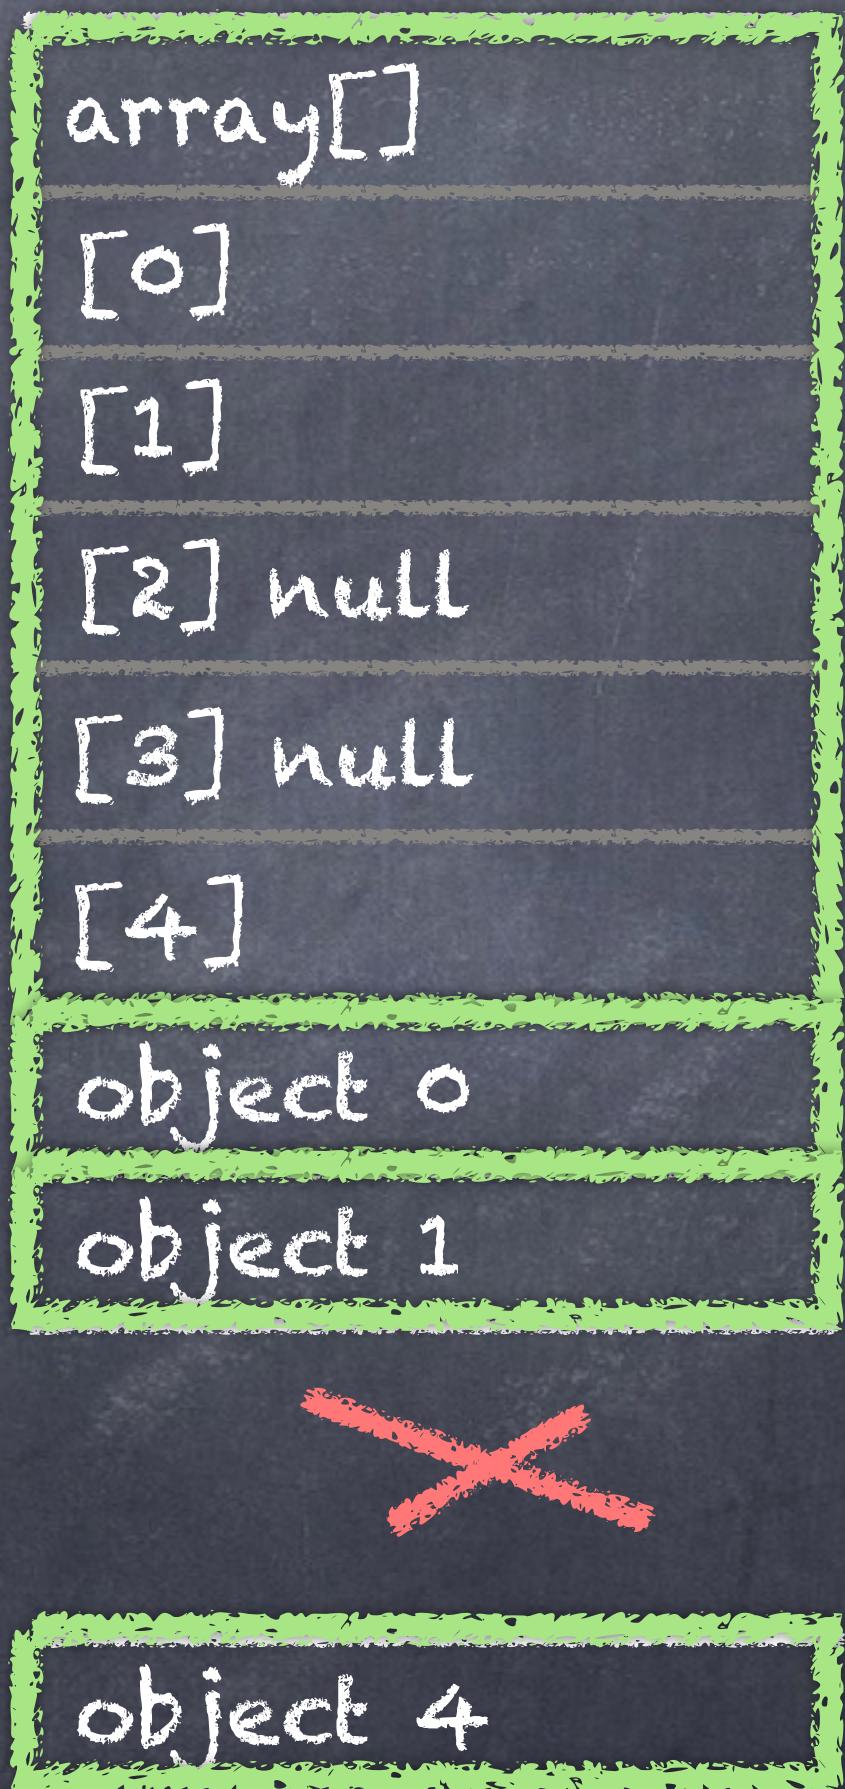
Garbage
collection runs:

All dead objects
are **swept**

Stack



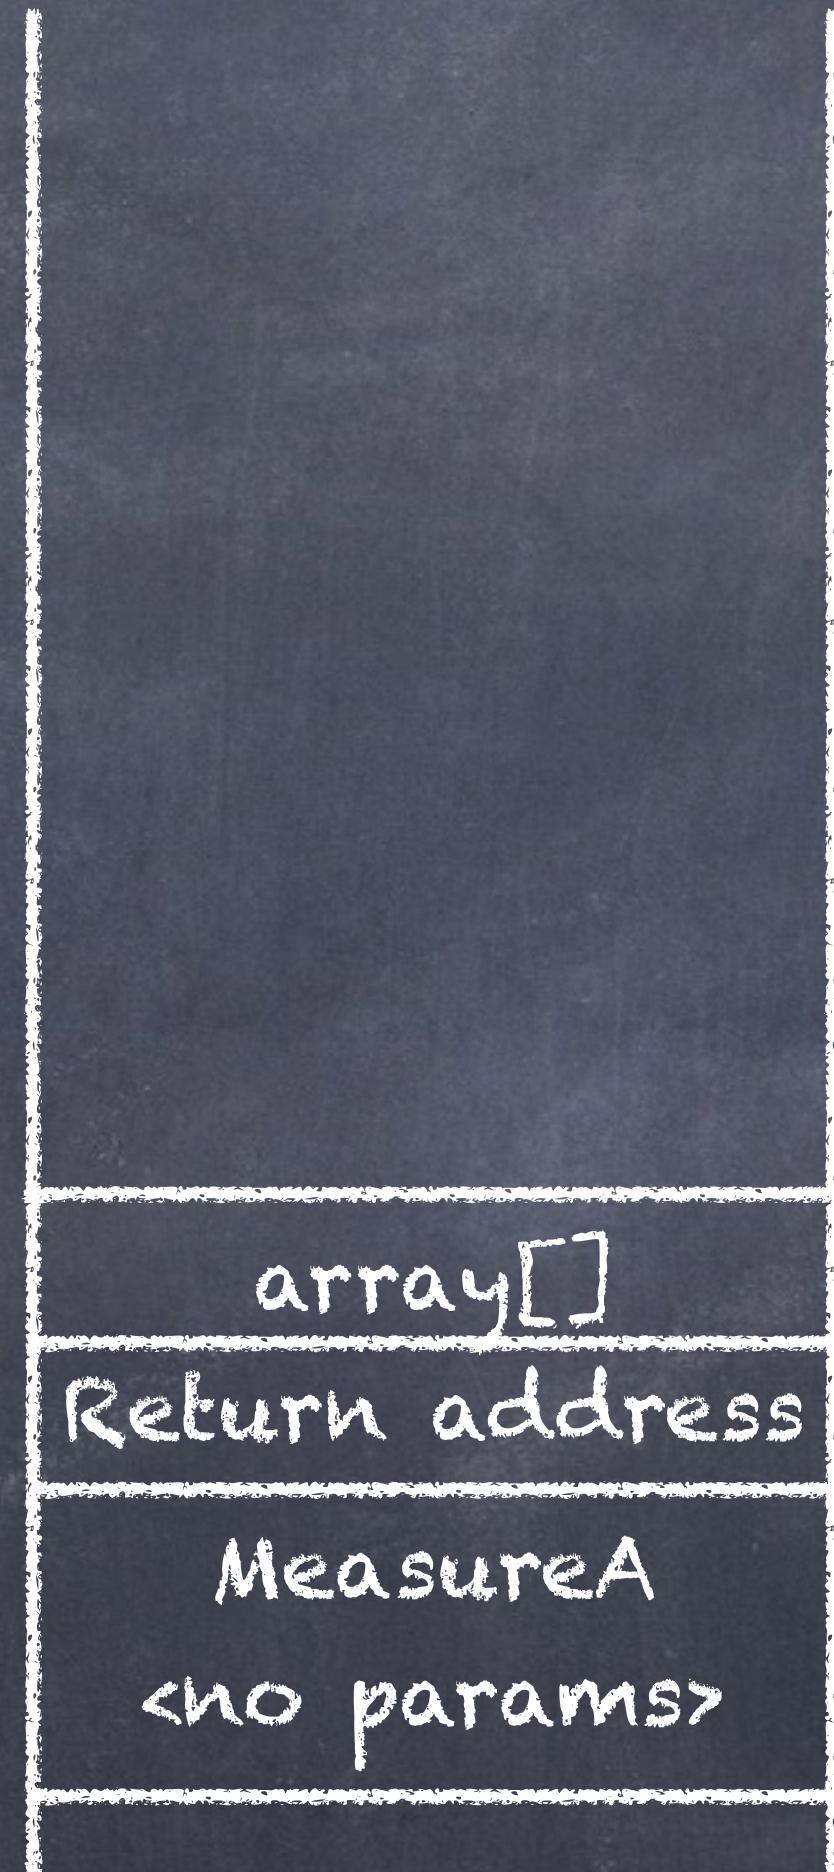
Heap



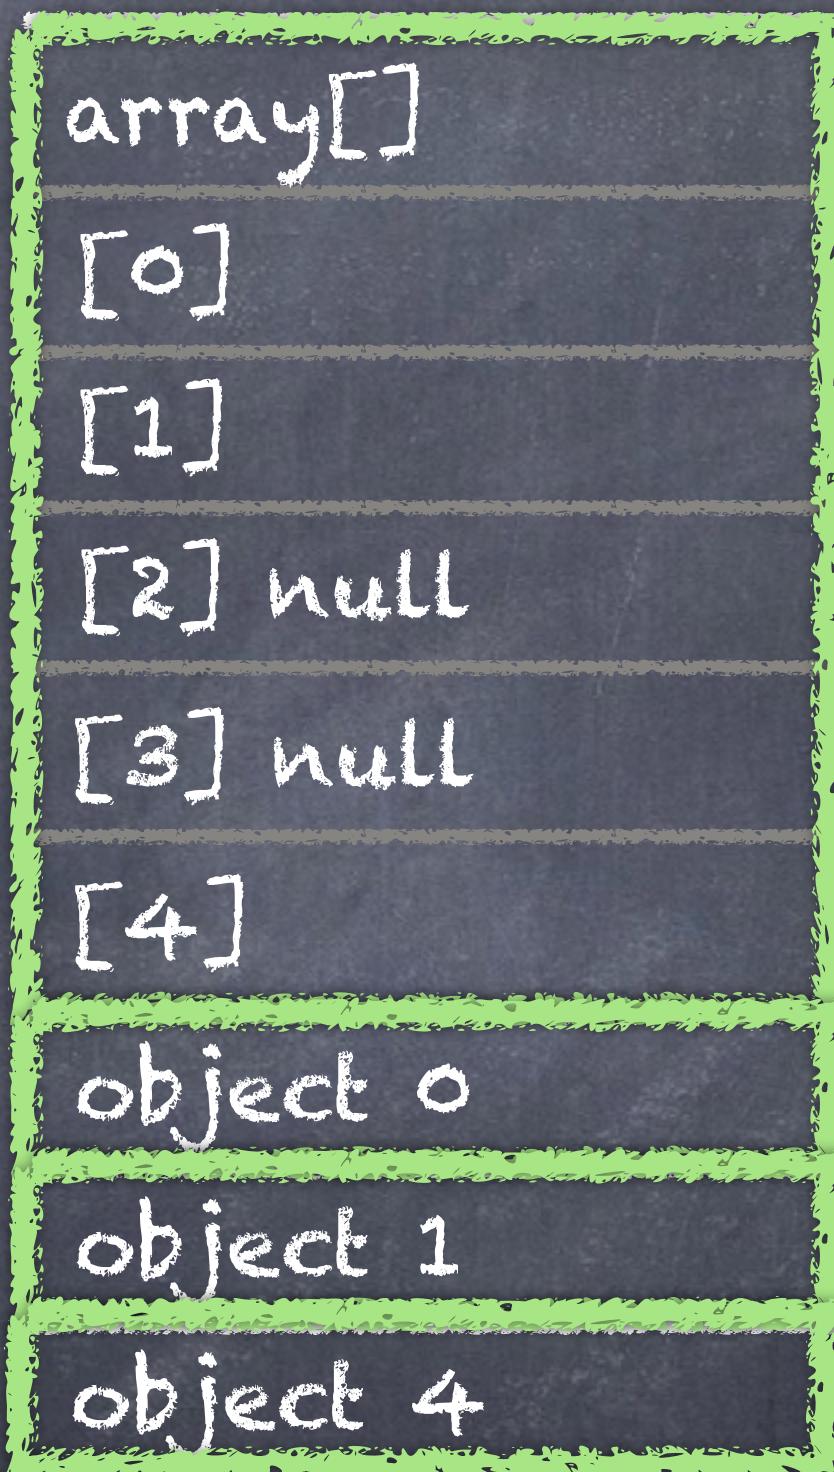
Garbage
collection runs:

The heap is
compacted

Stack



Heap



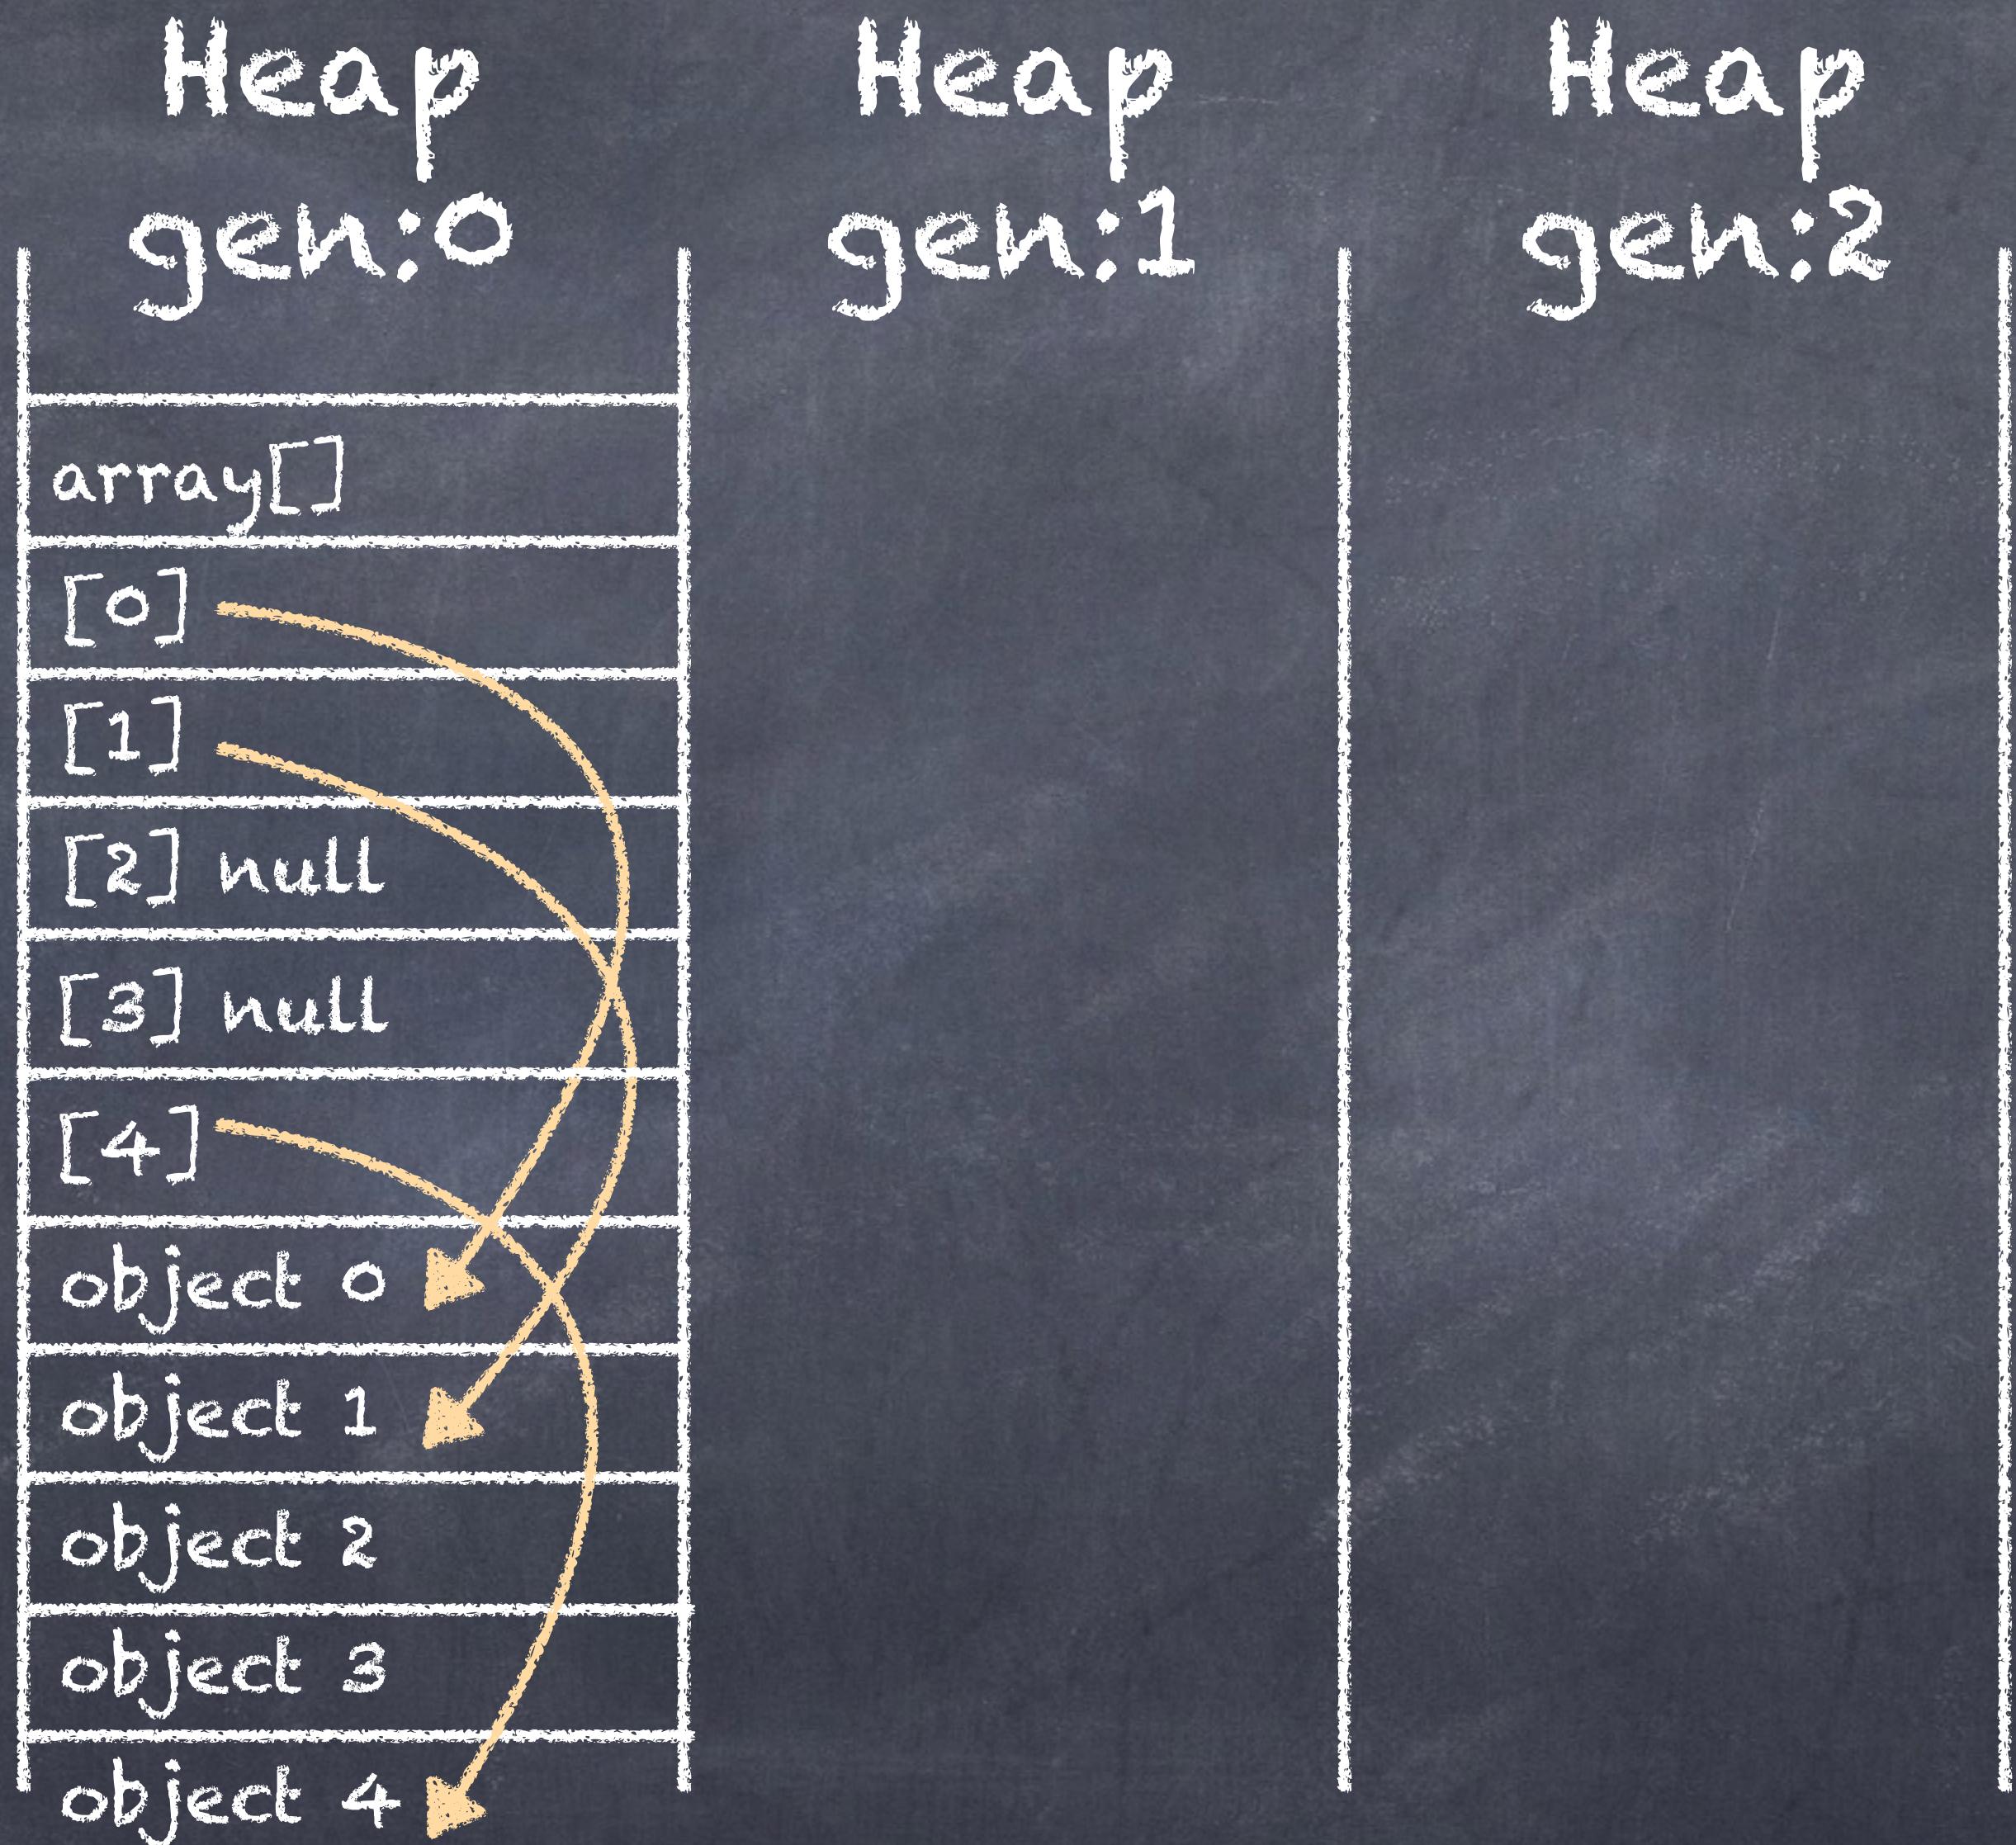
First problem

- A mark and sweep Garbage Collector has to mark all live objects on the heap during each collection cycle.
- This can lead to long program freezes while the collection is running
- Also: repeatedly marking the same objects over and over in each cycle is very inefficient

Solution: Generational Garbage Collector

Initial memory layout:

All objects are in gen 0



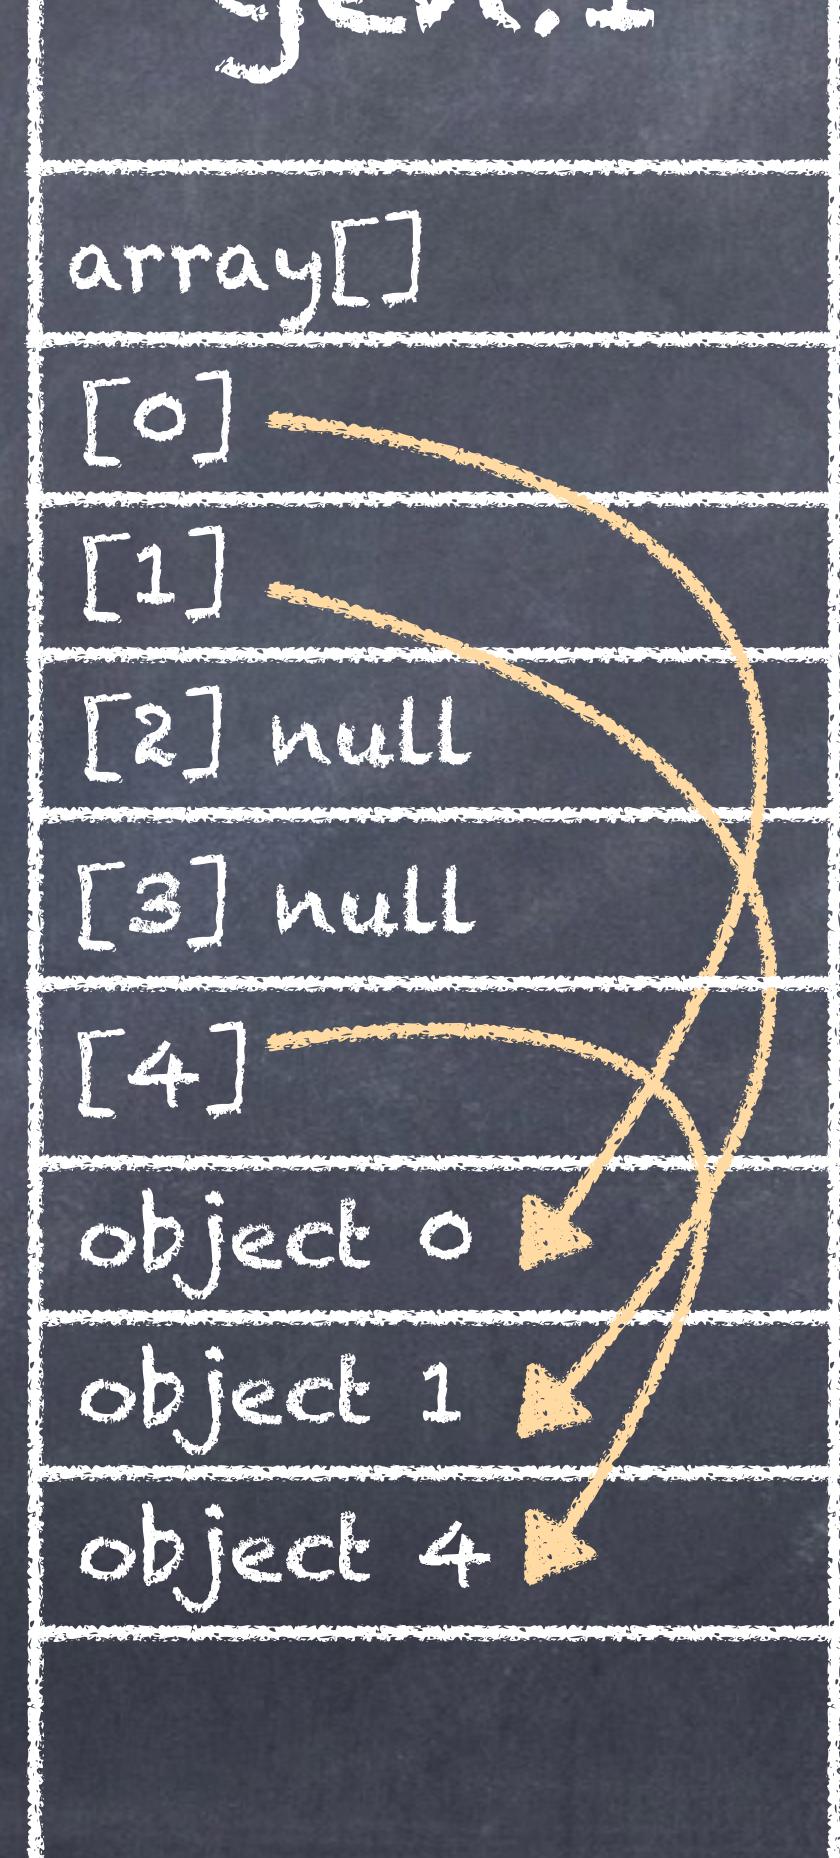
After the first
collection cycle:

All remaining
live objects
move to gen 1

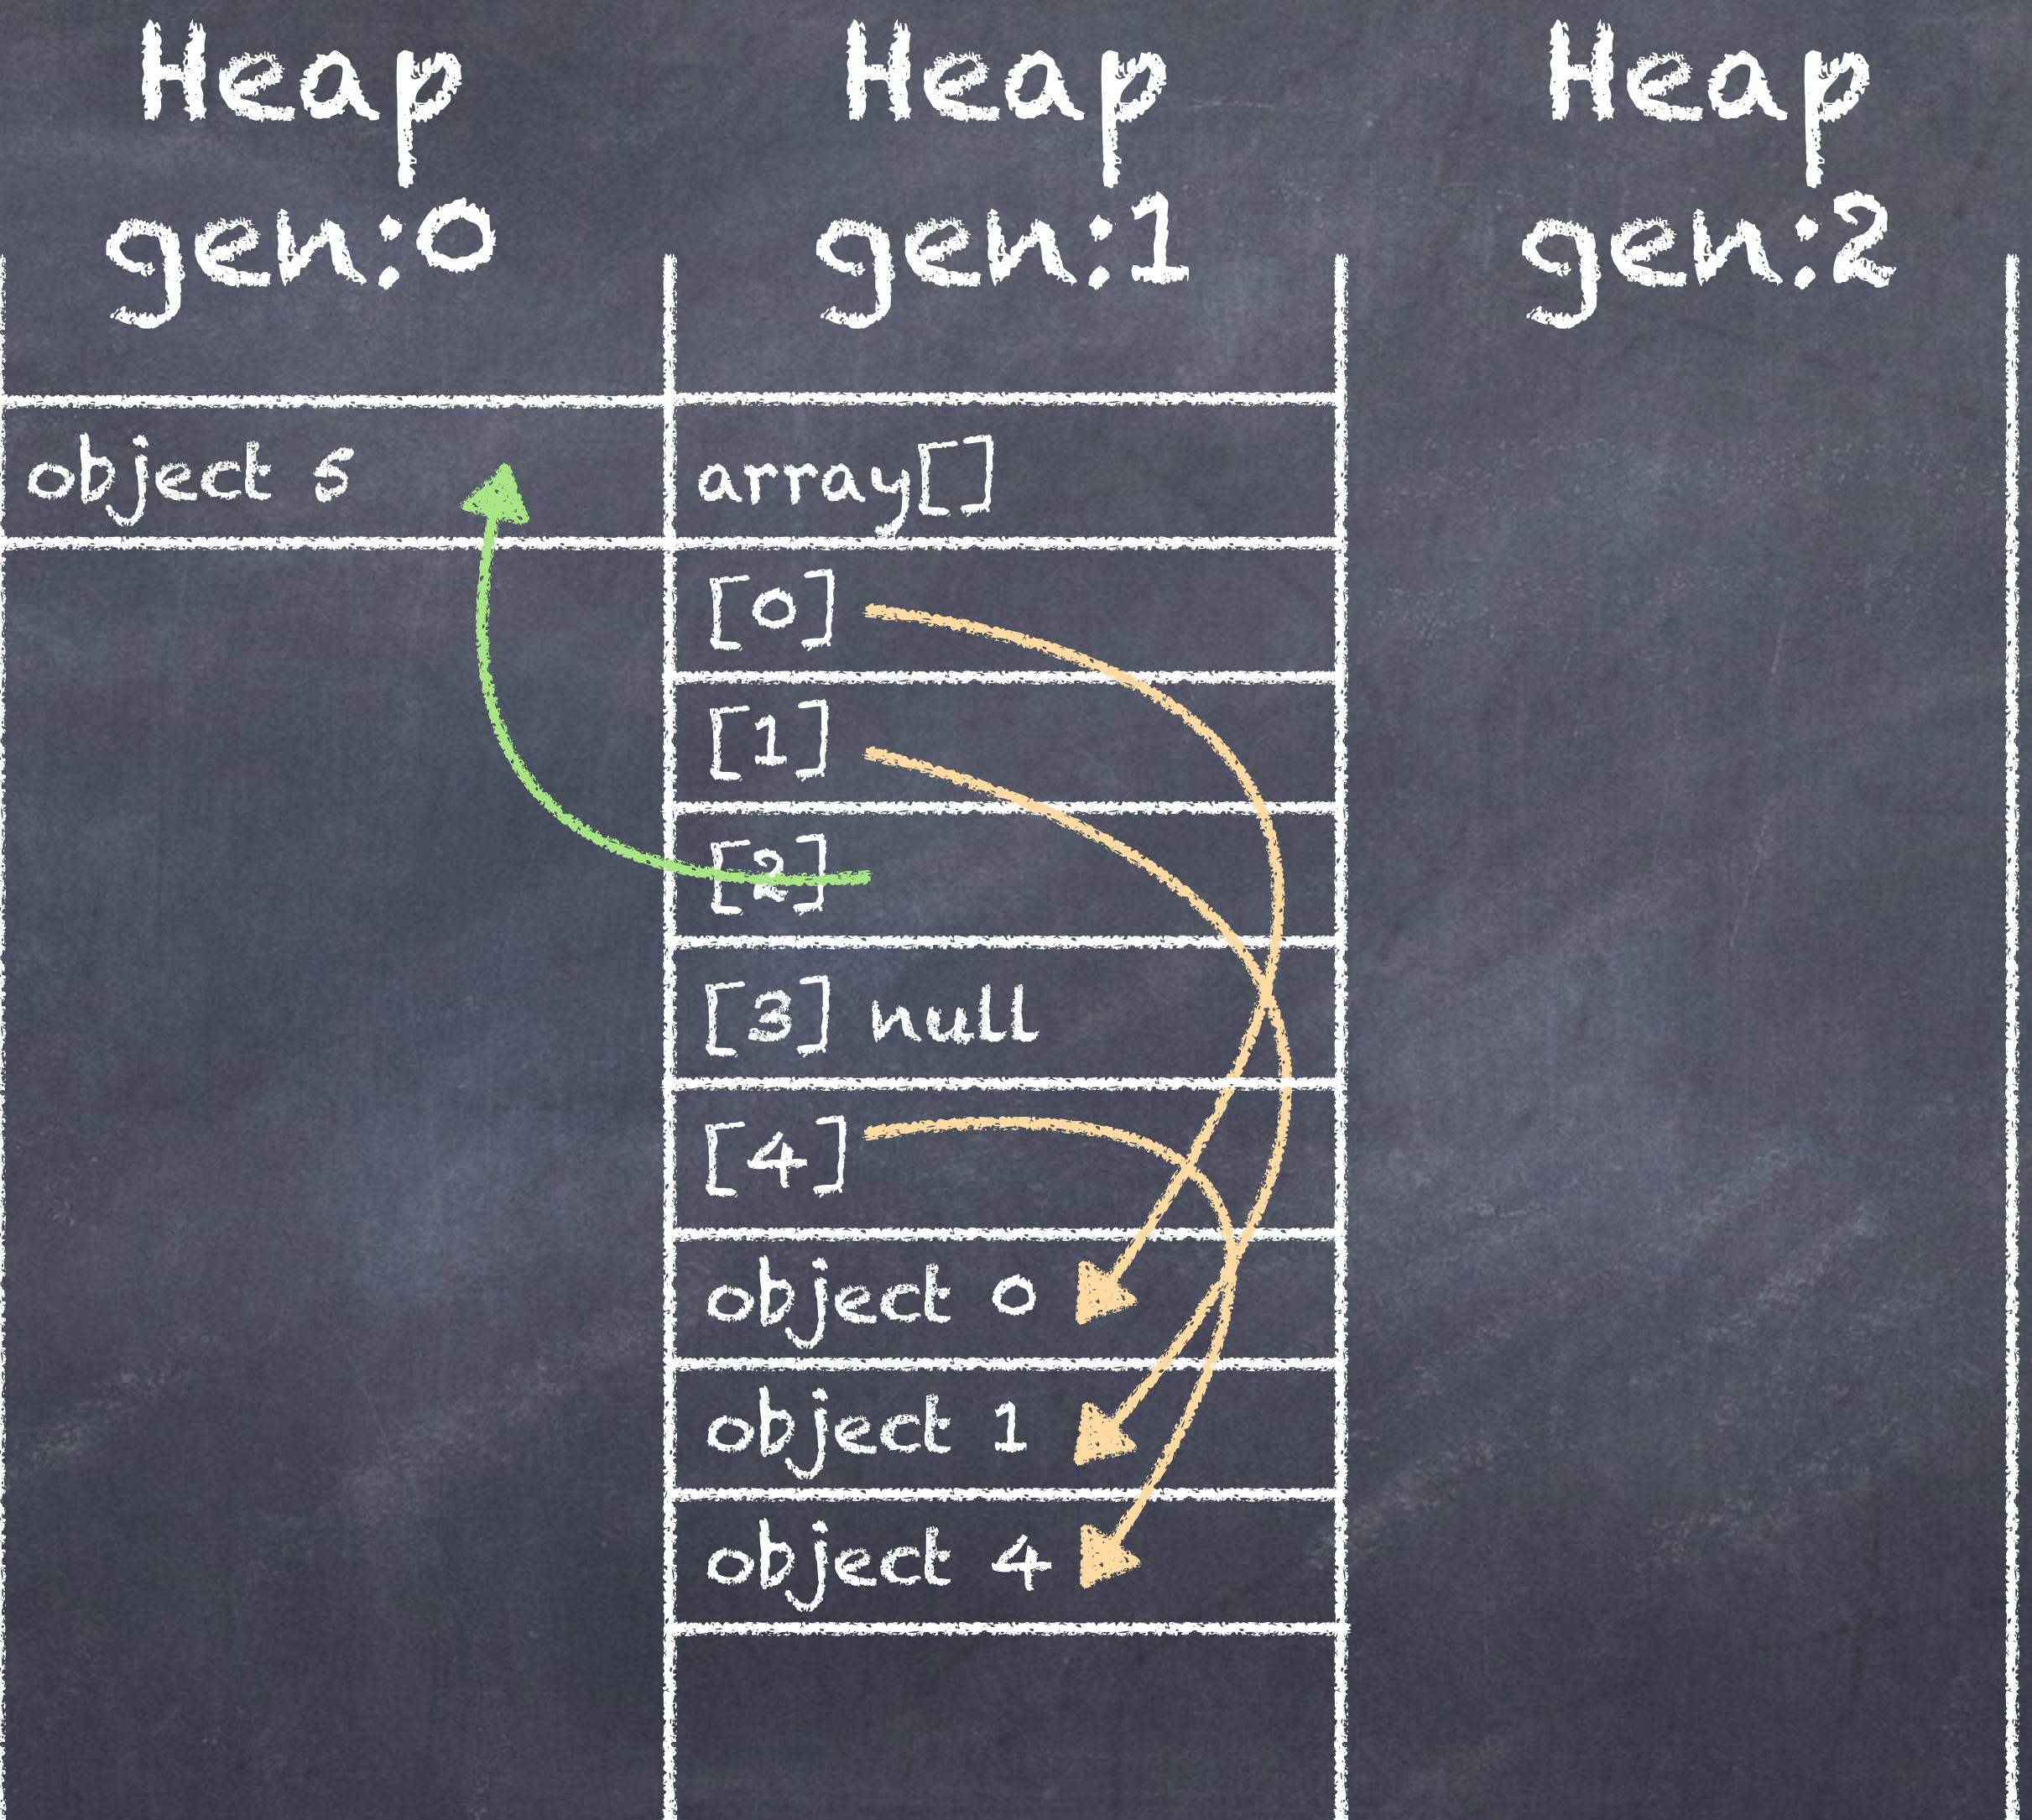
Heap
gen:0

Heap
gen:1

Heap
gen:2



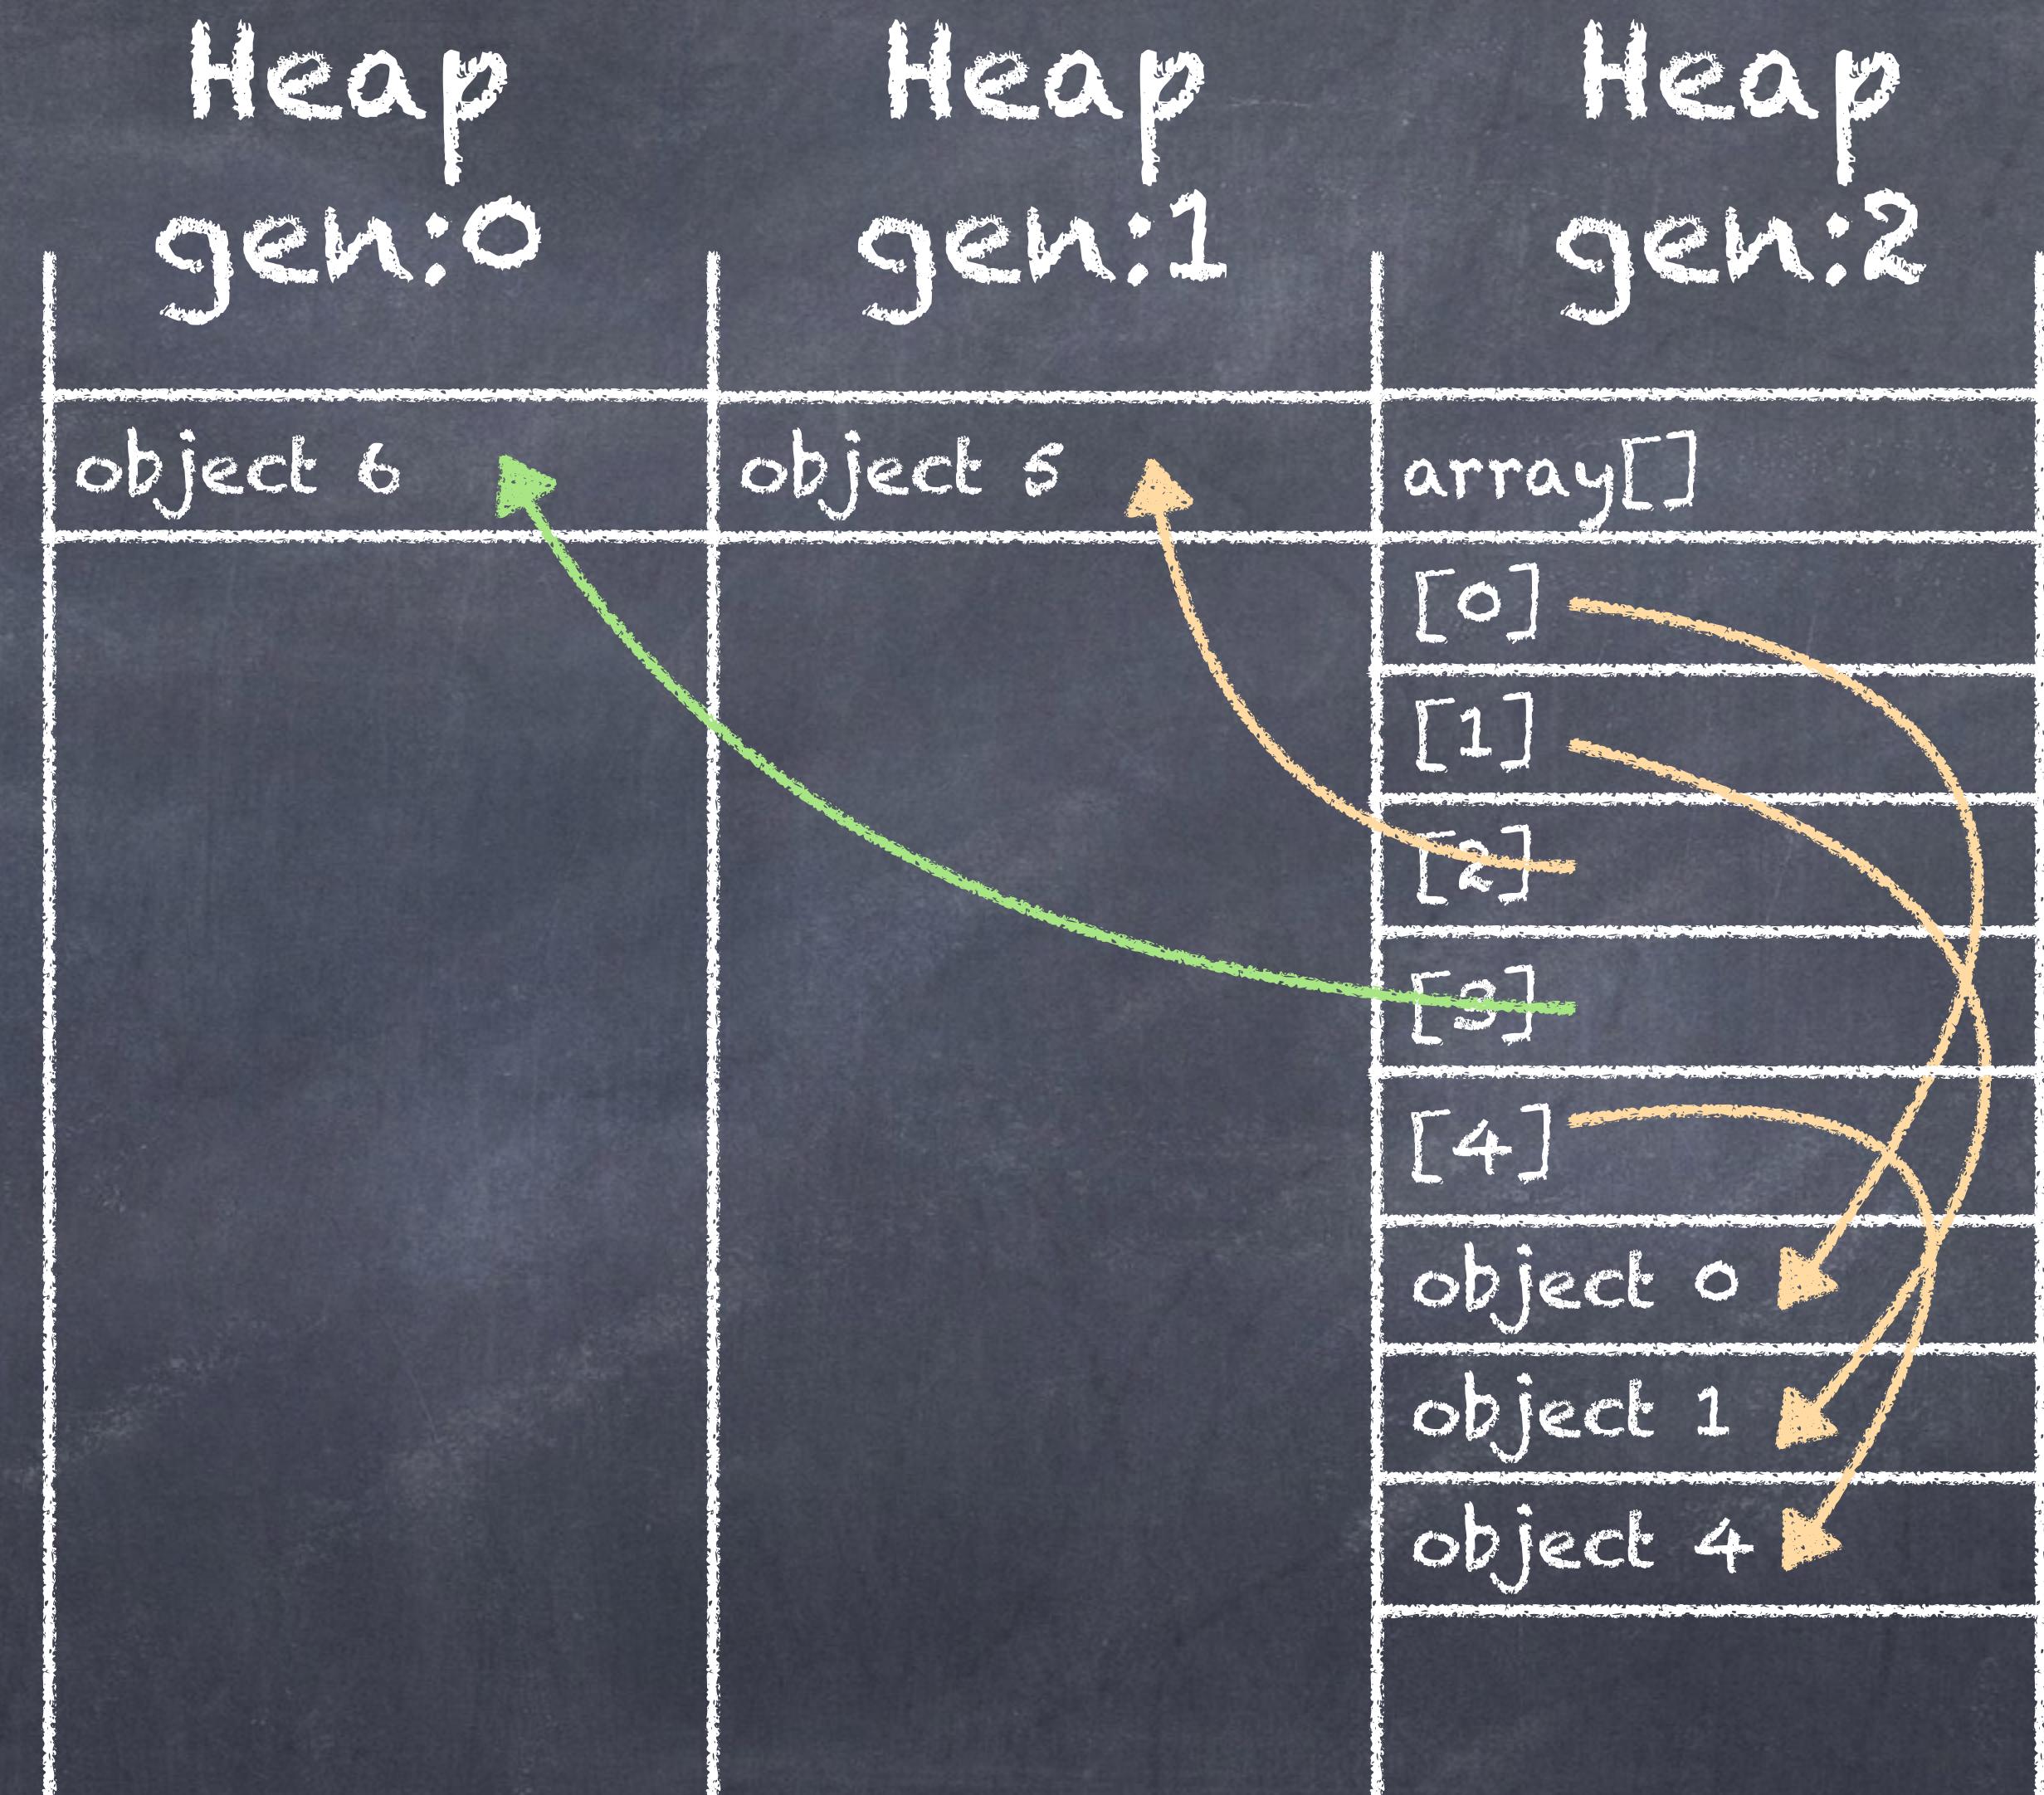
New allocations
go to gen 0



After the second collection cycle:

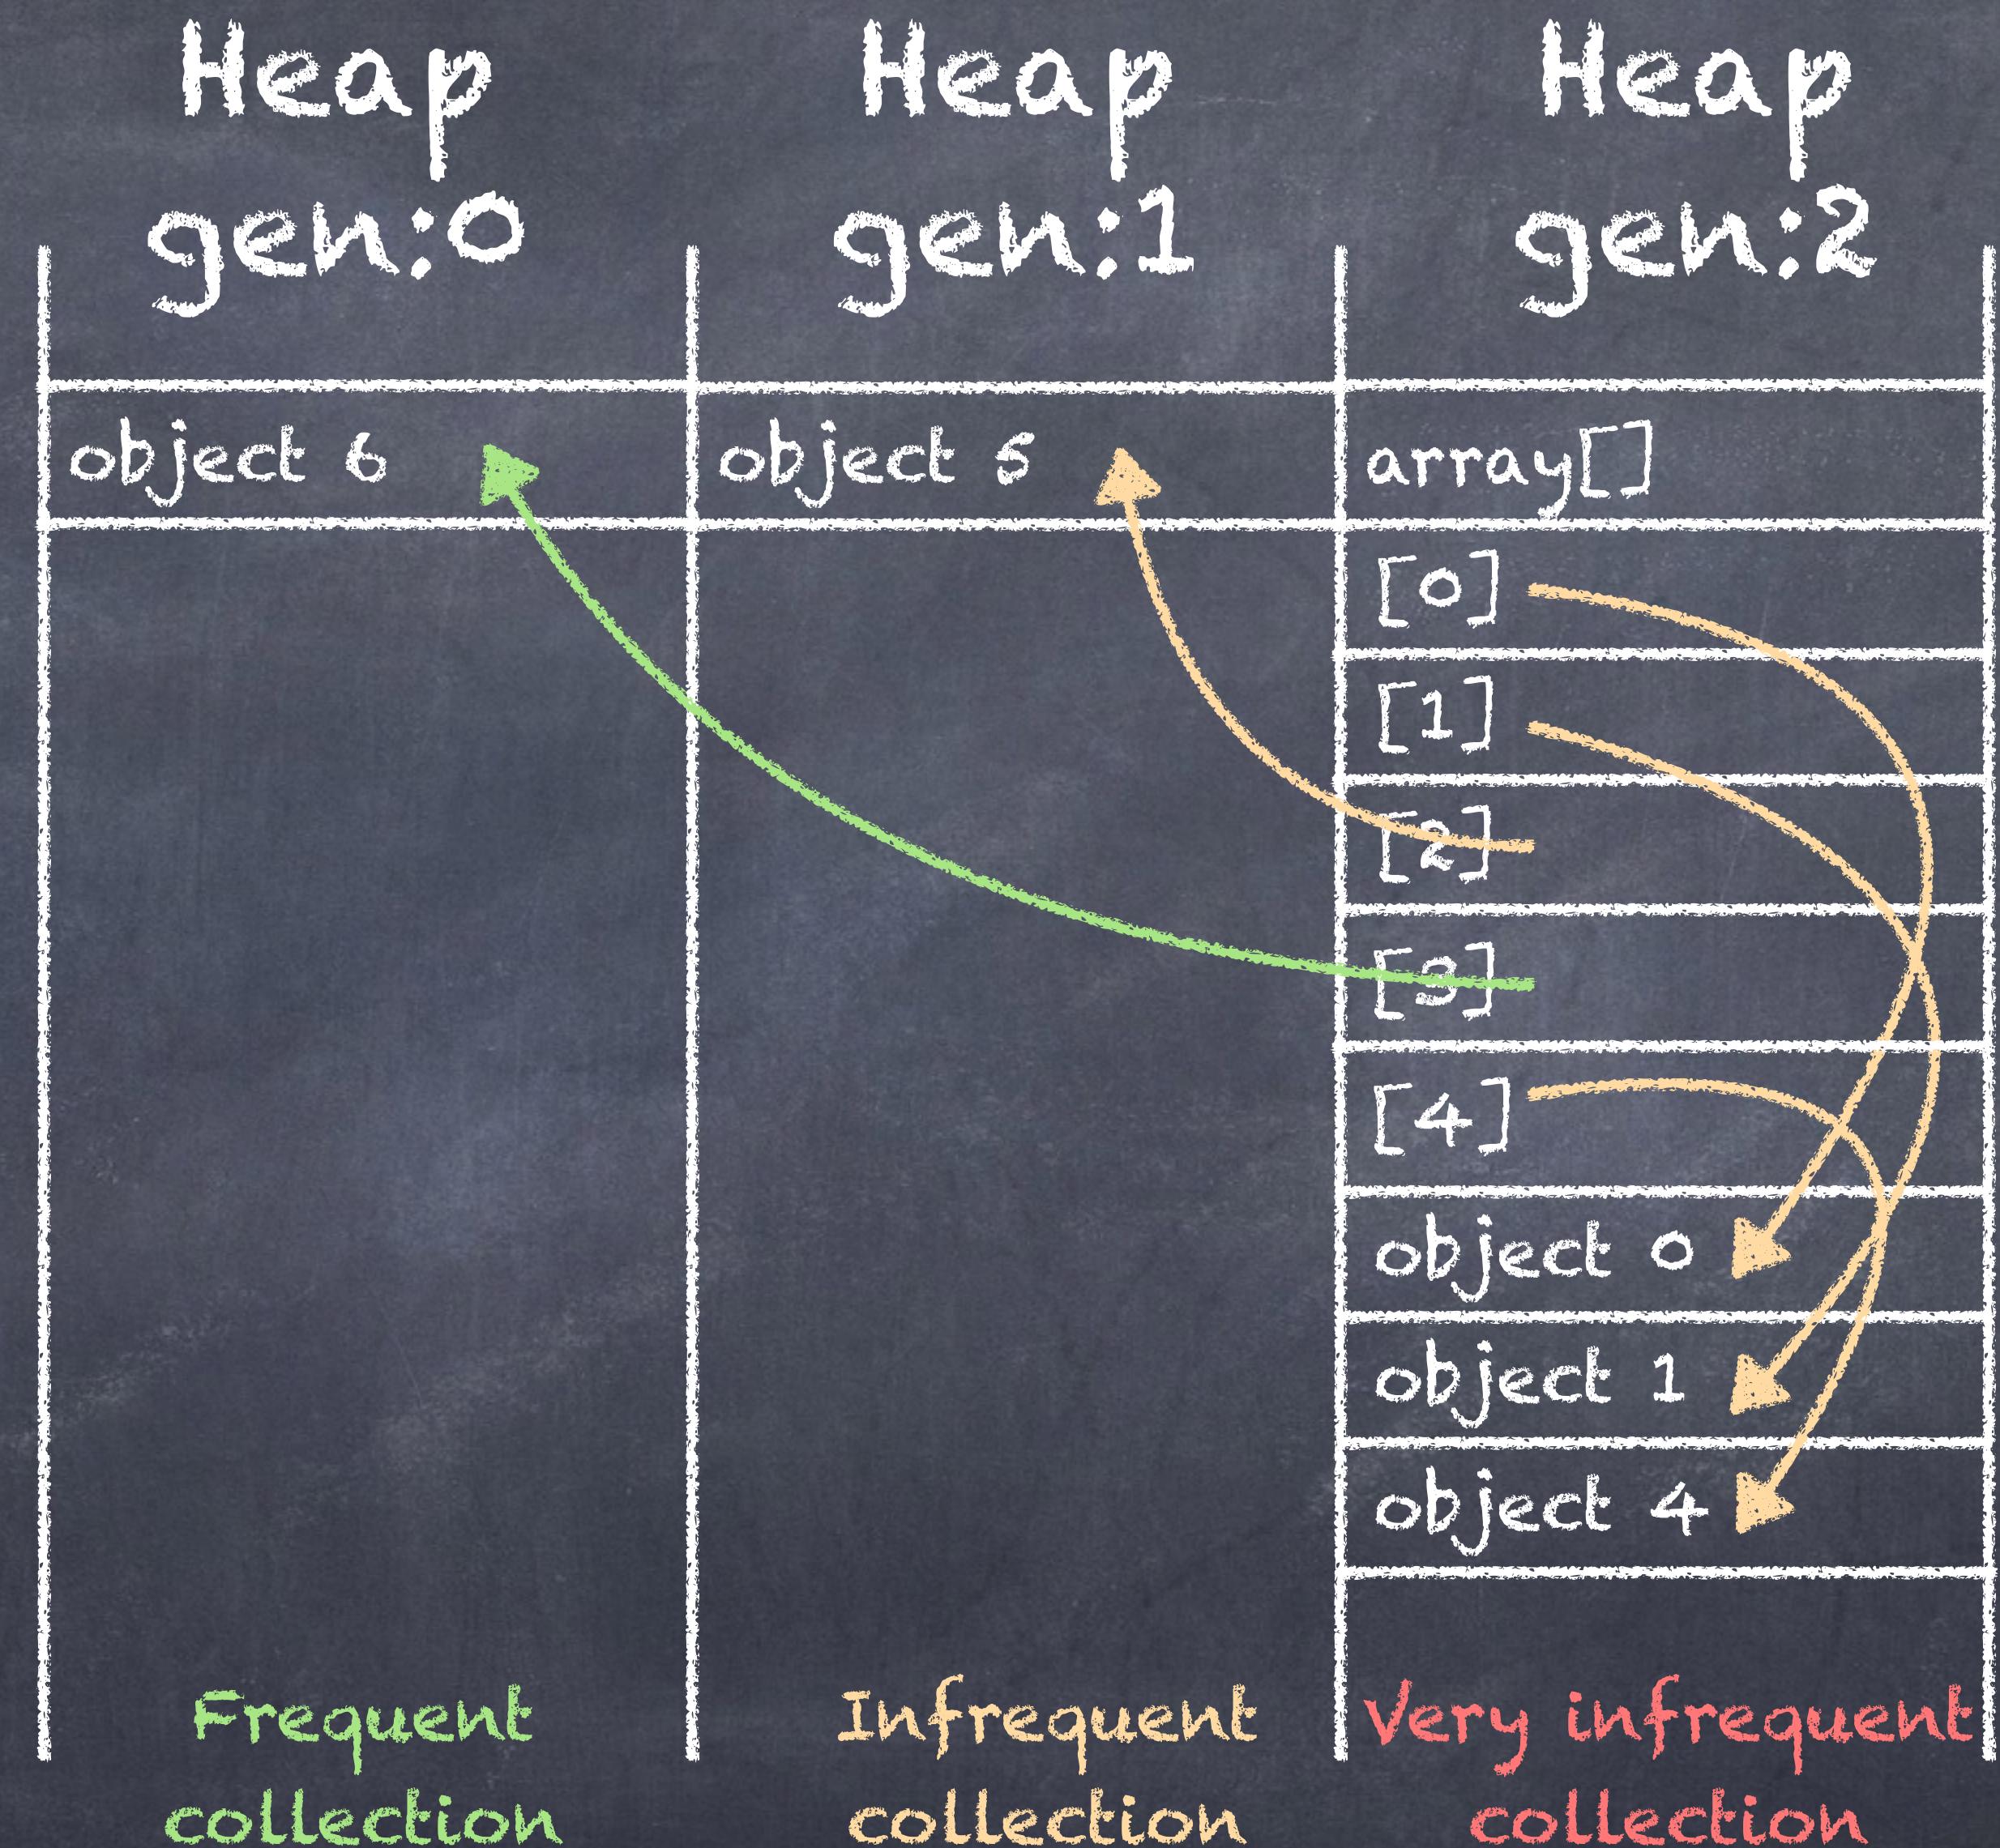
Gen 1 moves to gen 2 and gen 0 moves to gen 1

New allocations go to gen 0



Generations help
to reduce the
number of
objects in gen 0

Gen 1 and 2 are
collected much
less frequently



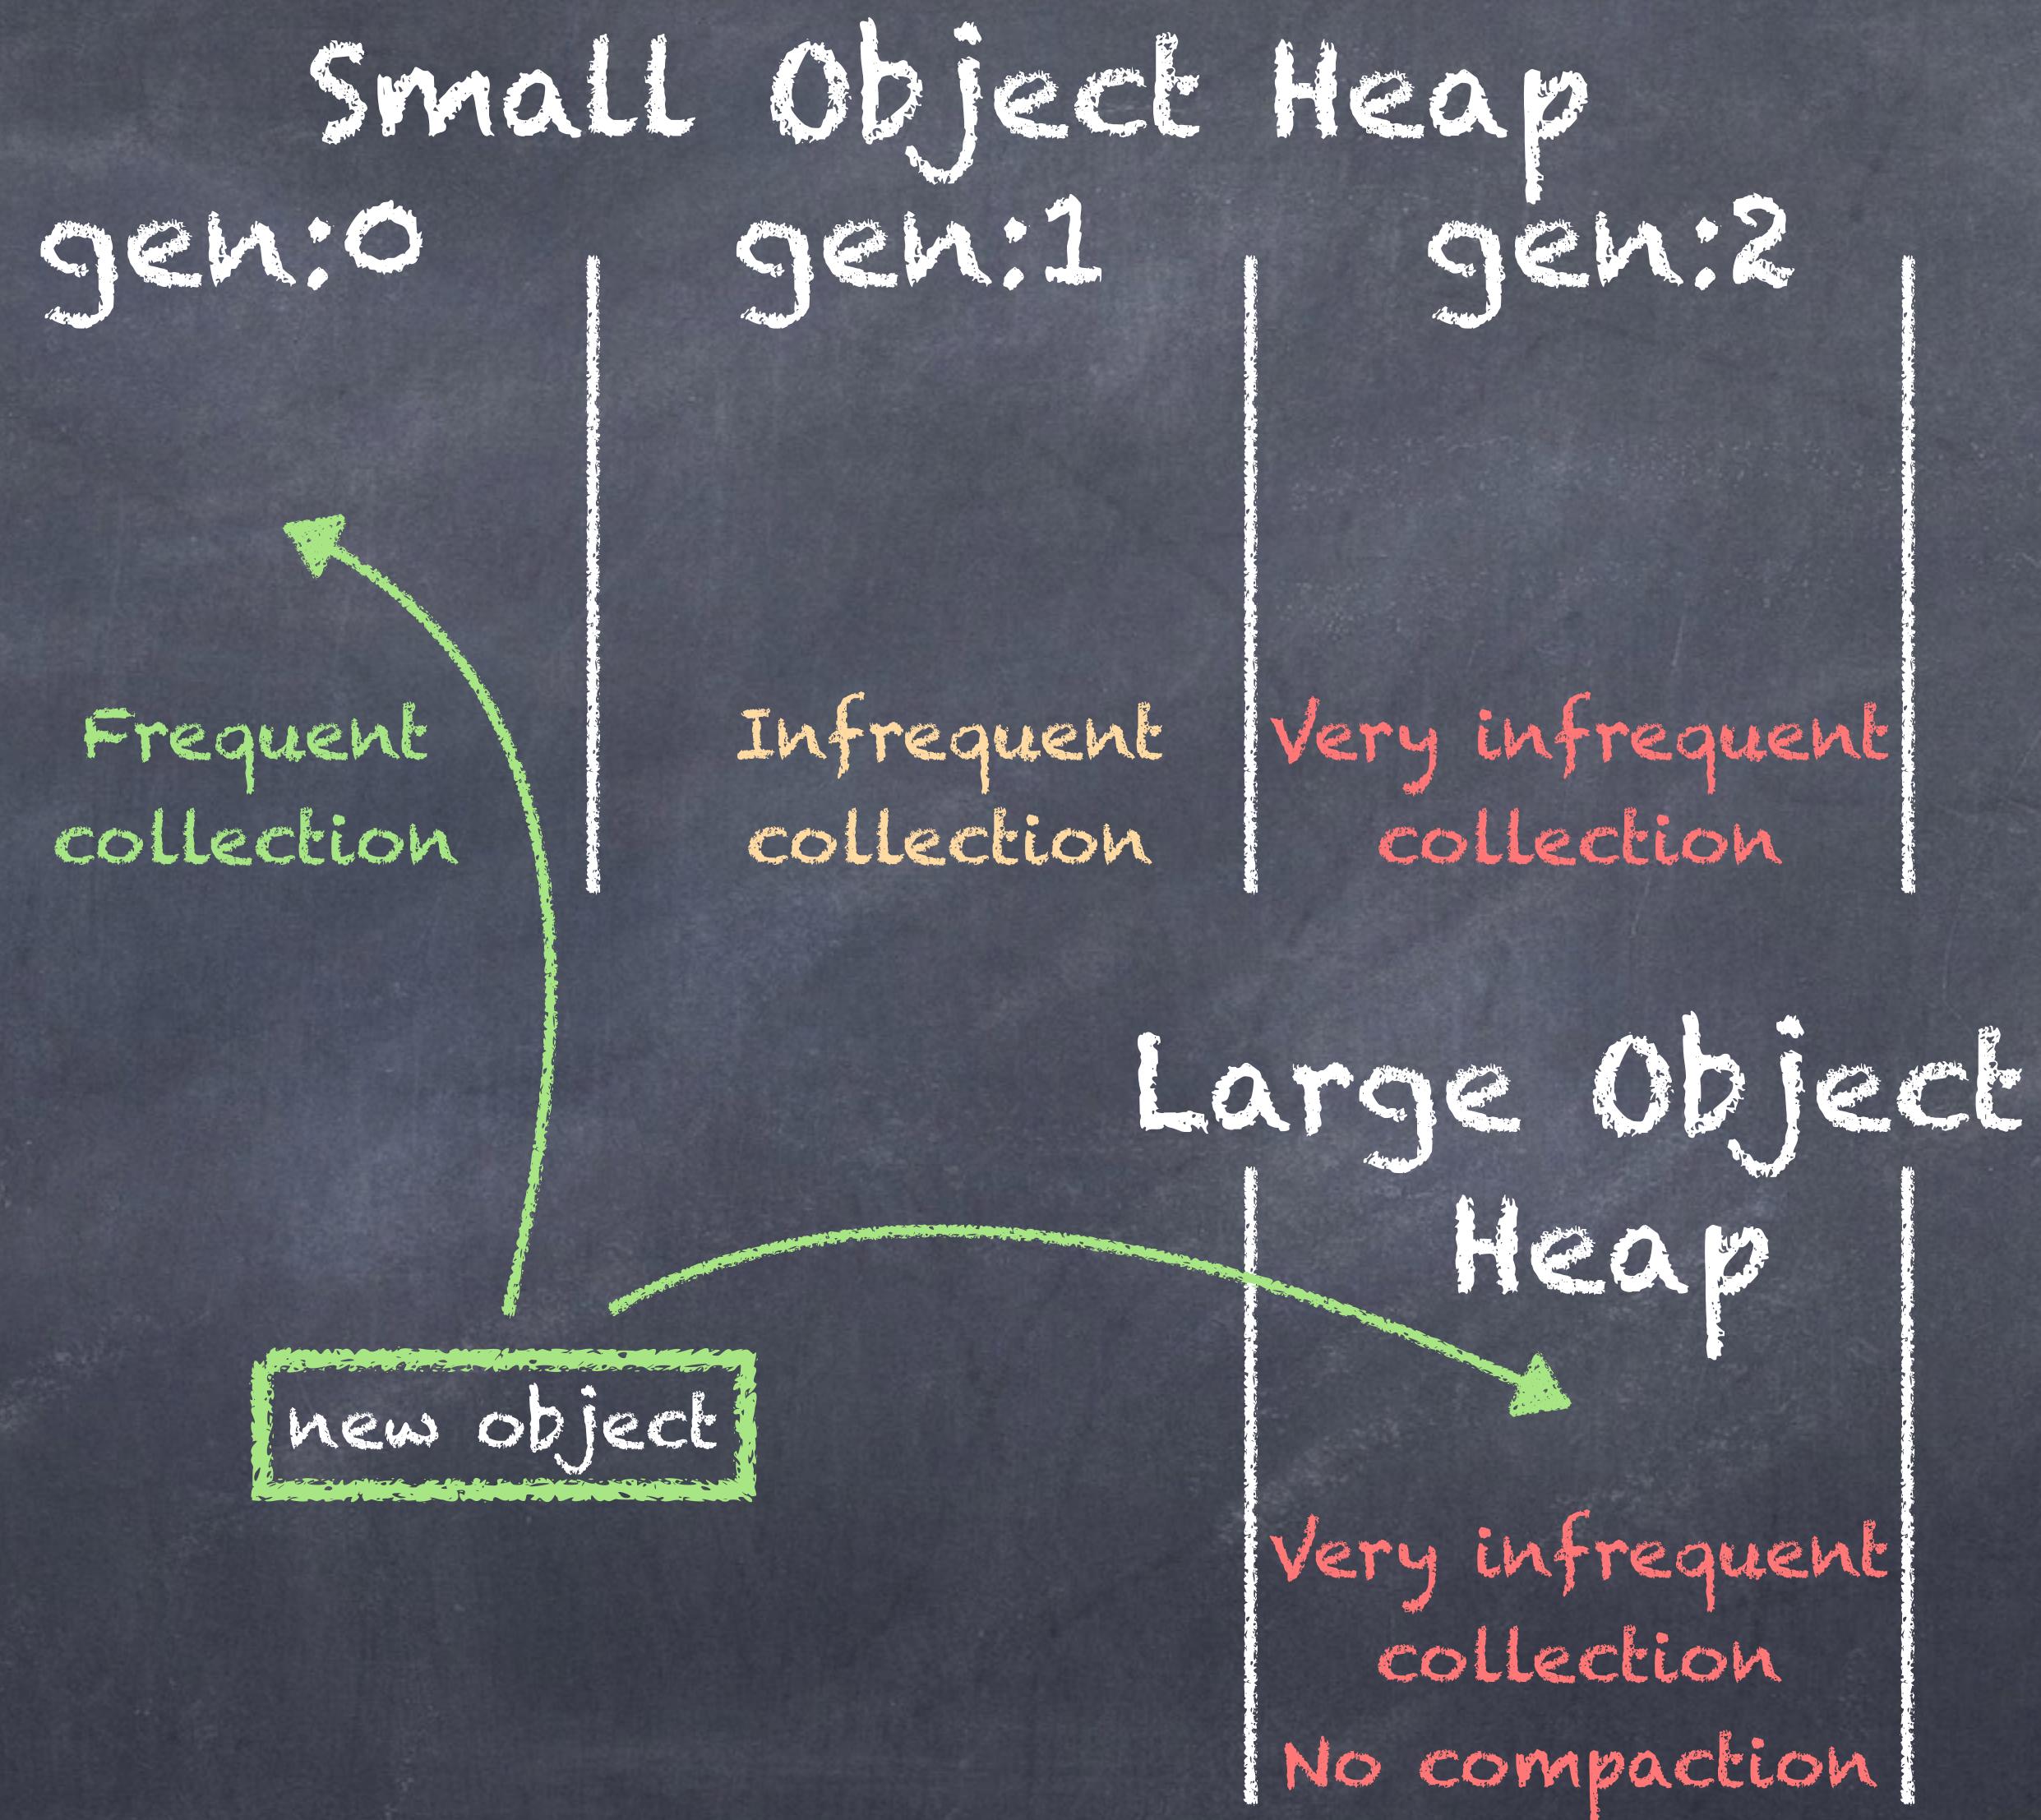
Second problem

- During 2 successive garbage collection cycles a long-living object will be compacted twice and moved twice:
4 memcpy operations for a single object
- For long-lived very large objects these 4 memory copy operations impact performance

Solution: Small Object and Large Object heaps

Objects larger than 85k are allocated on the Large Object Heap

The LOH is collected in gen 2 and not compacted



Implicit assumptions

- Objects are either short-lived or long-lived
- Short-lived objects will be allocated and discarded within a single collection cycle
- Objects that survive two collection cycles are long-lived
- 90% of all small objects are short-lived
- All large objects ($85K+$) are long-lived
- Do not go against these assumptions

What have we learned

- The Garbage Collector uses a mark/sweep/compact cycle
- Small objects are allocated on the Small Object Heap (SOH), large objects on the Large Object Heap (LOH)
- The Small Object Heap (SOH) has 3 generations. Objects are allocated in gen 0 and progress towards gen 2.
- The LOH has only one gen and does not compact
- .NET assumes 90% of all small objects are short-lived
- .NET assumes all large objects are long-lived